

Uncomputable Fractal Dimensions, Kolmogorov Complexity Cascades, and the Chaitin Barrier in Physical Systems

Ruqing Chen

GUT Geoservice Inc., Montreal, Quebec, Canada

ruqing@hotmail.com

DOI: [10.5281/zenodo.23194419](https://doi.org/10.5281/zenodo.23194419)

Abstract

We study three questions at the interface of fractal geometry and algorithmic information theory.

First, we construct a computable compact set $\mathcal{K} \subset [0, 1]$ whose box-counting and Hausdorff dimensions both equal $\frac{1}{4} + \frac{1}{2}\Omega$, where Ω is Chaitin’s halting probability. The dimension is therefore left-computably-enumerable but not computable. The construction is a homogeneous Cantor set with stage-dependent ratios driven by computable lower approximations of Ω , and the dimension follows from a Cesàro argument and the mass distribution principle. We also check the limits of the construction: no self-similar set with finitely many computable ratios can have an uncomputable dimension, and the known non-computability of certain Julia sets concerns the sets rather than their dimensions.

Second, we classify the growth of the Kolmogorov complexity of ε -discretisations with $L = \ln(1/\varepsilon)$ into three tiers. It is $\Theta(\log L)$ for every computable set, including smooth manifolds, self-similar fractals and the set above; $\Theta(L)$ when one incompressible symbol is injected per scale; and $e^{\Theta(L)}$ for independent randomness in every cell. Every question about a finite discretisation is decidable, and undecidability enters only through limits, as in the undecidability of the spectral gap.

Third, we test quantitatively the idea that physical cutoffs act as a “firewall” protecting physics from Gödel incompleteness, and find that the numbers refute it. Independence from ZFC already occurs for a Turing machine whose description is of order 10^5 bits, while a macroscopic crystal at atomic resolution carries up to 10^{24} bits of configurational information, and even a one-dimensional random fractal crosses the barrier at $L^* \approx 34.5$, far above the Planck scale ($L_P \approx 80$ for a 1 m system).

Numerical illustrations use binary lambda calculus, whose closed-term counts reproduce OEIS A114852, and LZMA compression of four classes of Cantor-type sets.

Contents

1	Introduction	2
2	A computable set with uncomputable fractal dimension	2
2.1	Statement	2
2.2	Proof	3
2.3	Scope of the construction	3
3	Three-tier hierarchy of Kolmogorov complexity under scaling	4
4	Quantitative refutation of a “Planck firewall” against Gödel	4

5 Numerical illustrations	5
5.1 A toy halting probability in binary lambda calculus	5
5.2 Complexity and dimension of four Cantor-type sets	6
6 Discussion and physical scope	6
A Technical details	7

1 Introduction

Self-similarity is usually associated with simplicity: a finite iterated function system (IFS) describes a fractal with infinitely many details by a few numbers, and its dimension follows from a single equation. On the other hand, algorithmic information theory [1, 2] distinguishes real numbers and infinite objects that can be computed to arbitrary precision from those that cannot. The canonical uncomputable real is Chaitin’s halting probability [3]

$$\Omega = \sum_{p \in \text{dom}(U)} 2^{-|p|}, \quad (1)$$

where U is a universal prefix-free machine. Ω is left-computably-enumerable (left-c.e.): it is the limit of a computable non-decreasing sequence of rationals Ω_s . However, it has no computable modulus of convergence, and its binary expansion is Martin-Löf random [4, 3].

This paper asks three questions.

1. Can a set that a computer draws to any desired precision have an uncomputable fractal dimension?
2. How does the Kolmogorov complexity of an ε -discretised geometric object grow as $\varepsilon \rightarrow 0$, and where, if anywhere, does undecidability enter?
3. Are physical cutoffs, such as the Planck length or a lattice spacing, needed to protect physical laws from Gödel incompleteness and the halting problem?

Our answers are: yes, with an explicit construction and proof (Section 2); a three-tier classification, with undecidability confined to limits (Section 3); and no, by a quantitative argument (Section 4). Section 5 presents numerical illustrations, and Section 6 discusses scope and limitations.

Related work. Non-computable quadratic Julia sets were constructed by Braverman and Yampolsky [5]. Undecidability of a physical property, the spectral gap of a translationally invariant lattice Hamiltonian, was proved by Cubitt, Pérez-García and Wolf [6], and Moore showed that low-dimensional dynamical systems can be computationally universal, which makes long-time questions undecidable [7]. Chaitin’s information-theoretic incompleteness theorem bounds the complexity that a formal system can certify [8], and Yedidia and Aaronson gave an explicit 7910-state Turing machine whose behaviour is independent of ZFC [9]. For concrete prefix-free machines see Tromp’s binary lambda calculus [10] and Stay [11].

2 A computable set with uncomputable fractal dimension

2.1 Statement

A compact set $\mathcal{K} \subset \mathbb{R}$ is *computable* if there is an algorithm that, given k , outputs a finite union of rational intervals whose Hausdorff distance to $\mathcal{K}$ is at most 2^{-k} [12].

Theorem 1. Let $(\Omega_s)_{s \geq 1}$ be a computable non-decreasing sequence of rationals in $[0, 1)$ converging to Ω . Put $d_k = \frac{1}{4} + \frac{1}{2}\Omega_k$ and $r_k = 2^{-1/d_k}$. Let $\mathcal{K} = \bigcap_k E_k$, where $E_0 = [0, 1]$ and E_k is obtained from E_{k-1} by replacing every interval I by its two end pieces of length $r_k|I|$. Then $\mathcal{K}$ is a computable compact set and

$$\dim_B \mathcal{K} = \dim_H \mathcal{K} = D \equiv \frac{1}{4} + \frac{1}{2}\Omega. \quad (2)$$

In particular $\dim_H \mathcal{K}$ is left-c.e. but not computable.

2.2 Proof

Since $\Omega_k \in [0, 1)$, we have $d_k \in [\frac{1}{4}, \frac{3}{4})$ and

$$r_k \in [2^{-4}, 2^{-4/3}), \quad 1 - 2r_k > 1 - 2^{-1/3} > 0.206. \quad (3)$$

Write $\delta_k = \prod_{j \leq k} r_j$, the common length of the 2^k intervals of E_k .

Computability. Each E_k is a finite union of intervals whose endpoints are computable reals, uniformly in k . Every interval of E_k contains points of $\mathcal{K}$, and $\mathcal{K} \subset E_k$, so the Hausdorff distance between E_k and $\mathcal{K}$ is at most $\delta_k \leq 2^{-4k/3}$. Rational approximations of the endpoints to within 2^{-k} therefore give the required approximations, and $\mathcal{K}$ is computable.

Box dimension. Two intervals of E_k with the same parent are separated by a gap of length $(1 - 2r_k)\delta_{k-1} \geq 0.206\delta_{k-1}$, and gaps created at earlier stages are larger. For $\delta_k \leq \delta < \delta_{k-1}$, each interval of E_k meets at most two δ -mesh cells, so $N_\delta(\mathcal{K}) \leq 2^{k+1}$. Conversely, the 2^{k-1} intervals of E_{k-1} are separated by gaps of at least $0.206\delta_{k-2} \geq 0.206\delta$, so each δ -cell meets at most a bounded number C of them and $N_\delta(\mathcal{K}) \geq 2^{k-1}/C$. Since $r_k \geq 2^{-4}$, $\ln(1/\delta) = \ln(1/\delta_k) + O(1)$. Hence

$$\dim_B \mathcal{K} = \lim_{k \rightarrow \infty} \frac{k \ln 2}{\ln(1/\delta_k)} = \lim_{k \rightarrow \infty} \frac{k \ln 2}{\sum_{j \leq k} \ln 2/d_j} = \lim_{k \rightarrow \infty} \left(\frac{1}{k} \sum_{j \leq k} \frac{1}{d_j} \right)^{-1} = D, \quad (4)$$

because $1/d_j \rightarrow 1/D$ and Cesàro means of a convergent sequence converge to the same limit.

Hausdorff dimension. Let μ be the natural measure, giving mass 2^{-k} to every interval of E_k . For a ball B of radius δ with $\delta_{k+1} \leq \delta < \delta_k$, the gaps between intervals of E_k are at least $0.206\delta_{k-1} \geq 0.206\delta_k$, so B meets at most $M = \lceil 2/0.206 \rceil + 2 = 12$ of them. Hence $\mu(B) \leq 12 \cdot 2^{-k}$. Given $\eta > 0$, for large k we have $2^{-k} \leq \delta_{k+1}^{D-\eta} \leq \delta^{D-\eta}$ by the box-dimension computation, so $\mu(B) \leq 12\delta^{D-\eta}$. The mass distribution principle [13] gives $\dim_H \mathcal{K} \geq D - \eta$ for every $\eta > 0$. Together with $\dim_H \leq \dim_B = D$, this proves $\dim_H \mathcal{K} = D$.

Uncomputability. If D were computable, so would be $\Omega = 2D - \frac{1}{2}$, contradicting the uncomputability of Ω [3]. D is left-c.e. because Ω is. $\square$

2.3 Scope of the construction

Proposition 2.1 (Self-similar sets have computable dimension). *Let $\mathcal{K}$ be the attractor of a finite IFS of similarities with computable ratios $r_1, \dots, r_N$ satisfying the open set condition. Then $\dim_H \mathcal{K} = \dim_B \mathcal{K}$ is computable.*

Proof. The dimension is the unique root s of $\sum_i r_i^s = 1$ [14, 13]. The function $s \mapsto \sum_i r_i^s$ is computable and strictly decreasing, so the root can be computed to any precision by bisection. $\square$

Uncomputability of the dimension therefore requires infinitely many independent data, as in Theorem 1, where the ratios change at every stage.

Remark (Julia sets). Braverman and Yampolsky proved that there exist quadratic Julia sets that are not computable [5]. That result concerns computability of the *set*. It is logically independent of Theorem 1, in which the set is computable but its dimension is not.

Remark (Information content). The set of Theorem 1 has small Kolmogorov complexity at every finite resolution (Proposition 3.1). The uncomputable information resides only in the limit $k \rightarrow \infty$. Since the convergence $\Omega_s \rightarrow \Omega$ has no computable modulus, no algorithm can certify the dimension to a prescribed precision from finite-resolution data.

3 Three-tier hierarchy of Kolmogorov complexity under scaling

Let S_ε be the binary image of a set on the ε -grid of $[0, 1]^n$, with $\varepsilon = 2^{-k}$, and let $L = \ln(1/\varepsilon) = k \ln 2$. Let $K(\cdot)$ denote prefix Kolmogorov complexity [1].

Proposition 3.1 (Tier 1: computable sets). *For every computable compact set $\mathcal{K}$ (and a fixed rule for discretising it),*

$$K(S_\varepsilon) \leq K(k) + O(1) \leq \log_2 L + 2 \log_2 \log_2 L + O(1), \quad (5)$$

and $K(S_\varepsilon) \geq K(k) - O(1)$, so $K(S_\varepsilon) = \Theta(\log L)$ for all but a sparse set of k . This applies to smooth manifolds, to self-similar fractals and to the set of Theorem 1.

Proof. A fixed program that computes $\mathcal{K}$, given k , outputs $S_{2^{-k}}$, which proves the upper bound. Conversely, k can be read off from the size of S_ε , so $K(k) \leq K(S_\varepsilon) + O(1)$, and $K(k) \geq \log_2 k$ for most k . $\square$

Proposition 3.2 (Tier 2: one incompressible symbol per scale). *Let the ratio at stage j of a homogeneous Cantor construction be chosen from two fixed values according to the j -th bit x_j of an infinite binary sequence x . If x is Martin-Löf random (for instance $x = \Omega$), then $K(S_\varepsilon) = k \Theta(1) + O(\log k)$, i.e. $\Theta(L)$.*

Proof. The first k bits of x can be recovered from the stage- k image, because the stage- j ratio is read off from the gap structure. Hence $K(x \upharpoonright k) \leq K(S) + O(\log k)$. Martin-Löf randomness gives $K(x \upharpoonright k) \geq k - O(1)$ [3, 2]. The upper bound $K(S) \leq k + O(\log k)$ follows from describing $x \upharpoonright k$ literally. The resolution at stage k satisfies $L = \Theta(k)$. $\square$

Such a set is not computable, because it computes x .

Proposition 3.3 (Tier 3: independent randomness per cell). *If every interval chooses its own ratio by an independent fair bit, then with probability at least $1 - 2^{-c}$ one has $K(S_\varepsilon) \geq N_\varepsilon - c$, where $N_\varepsilon \sim \varepsilon^{-D} = e^{DL}$ is the number of decision bits up to resolution ε .*

Proof. The map from decision bits to images is injective, and at most 2^{N-c} strings of length N have complexity below $N - c$. $\square$

Decidability at finite resolution. Every question about a fixed S_ε is a question about a finite object and is therefore decidable. Complexity classes such as NP-hardness describe how cost scales with the size of the input. They do not describe a threshold resolution ε^* below which a problem suddenly becomes hard. Undecidability enters only through questions that quantify over all scales or all sizes. In the spectral-gap construction of [6], whether the system is gapped depends on whether a universal Turing machine halts, and the finite-size behaviour can mimic the opposite phase up to system sizes that admit no computable bound. The precise form of the “critical scale” intuition is therefore: *for each instance a crossover scale exists, but it is not a computable function of the instance.*

4 Quantitative refutation of a “Planck firewall” against Gödel

One may conjecture that a discrete physical cutoff, Planck or atomic, protects physical laws from Gödel incompleteness by halting the growth of algorithmic complexity before it reaches Chaitin’s barrier. We argue that this conjecture fails on four counts.

1. Logic. Incompleteness and the halting problem limit what a formal system can prove and what an algorithm can decide. They do not render physical laws inconsistent. A physical state whose properties are unprovable in a given theory is not a paradox.

2. Cutoffs do not remove undecidability. The spectral-gap problem is undecidable for Hamiltonians defined *on a lattice*, i.e. with a cutoff already in place [6]. The undecidability comes from the thermodynamic limit. Moore’s universal dynamical systems with three degrees of freedom are undecidable through the infinite-time limit [7]. Removing short distances leaves both limits intact.

3. The barrier is crossed long before any cutoff. By Chaitin’s incompleteness theorem, a formal system F cannot prove $K(x) > c_F$ for any specific x , where c_F is of the order of the complexity of F [8]. Yedidia and Aaronson exhibited a 7910-state, two-symbol Turing machine whose halting is independent of ZFC [9]. Encoding each of its 2×7910 transitions with about 15 bits gives an illustrative scale

$$c_F \approx 2.4 \times 10^5 \text{ bits.} \quad (6)$$

By comparison:

- a 1 cm^3 crystal at atomic resolution ($1 \text{ \AA}$) has 10^{24} sites, so its configurational information can reach 10^{24} bits;
- even the one-dimensional tier-3 random Cantor set of Section 5, with $D \approx 0.36$, reaches $K \sim e^{DL} = c_F$ at $L^* = \ln c_F / D \approx 34.5$, i.e. at $\varepsilon^* \approx 10^{-15}$ of the system size;
- the Planck cutoff for a 1 m system corresponds to $L_P = \ln(1 \text{ m} / 1.6 \times 10^{-35} \text{ m}) \approx 80$.

The premise that cutoffs come before the barrier is therefore false by many orders of magnitude.

4. The barrier has no dynamical content. Exceeding c_F changes nothing in the evolution of a system. It only limits which statements about its description F can prove.

A defensible weaker statement. If both an ultraviolet and an infrared cutoff hold, so that the relevant Hilbert space is finite-dimensional, then every question about a finite region over a finite time is decidable in principle. Undecidability in physics then arises only from idealised infinite limits. This follows from finiteness of the information content and has no relation to Gödel’s theorem.

5 Numerical illustrations

5.1 A toy halting probability in binary lambda calculus

In binary lambda calculus (BLC) [10], a term in de Bruijn notation is encoded as $00M$ for an abstraction λM , $01MN$ for an application MN , and $1^{i+1}0$ for the variable with index i . The encoding of closed terms is self-delimiting, so the set of closed-term codes is prefix-free. We define

$$\Omega_{\text{BLC}} = \sum_{\text{closed } M \text{ with a } \beta\text{-normal form}} 2^{-|M|}, \quad (7)$$

which is a left-c.e. real. We do not claim that it is Martin-Löf random, which would require an optimal universal prefix-free machine.

Our enumeration produces the numbers of closed terms of length $n = 4, 5, \dots, 24$: 1, 0, 1, 1, 2, 1, 6, 5, 13, 14, 37, 44, 101, 134, 298, 431, 883, 1361, 2736, 4405, 8574. This agrees term by term with OEIS A114852 [15], which independently checks the enumerator. Normal-order

reduction with a budget of 400 β -steps and 3000 nodes certifies normalisation for all but 34 of the 19 048 terms of length at most 24, namely 1, 1, 7 and 25 undecided terms at lengths 18, 20, 22 and 24. This gives the rigorous lower bound

$$\Omega_{\text{BLC}} \geq 0.1201817393\dots \quad (8)$$

Why no bits are certified. An upper bound requires bounding (a) the Kraft mass of all closed terms longer than the enumeration, and (b) the mass of enumerated terms that never normalise. Our computation provides neither. Without an analytic bound on the total closed-term measure, the remainder (a) can only be bounded by $1 - \sum_{|M| \leq 24} 2^{-|M|}$. Without proofs of non-termination, the undecided terms in (b) must be counted as possibly halting. The resulting upper bound is then trivially 1, and no leading bit of Ω_{BLC} is certified. This is the general situation. More computation improves lower bounds, but each further certified bit of a halting probability requires *proofs* that specific programs do not halt, and Chaitin’s theorem limits how many such proofs any fixed formal system can supply [8].

5.2 Complexity and dimension of four Cantor-type sets

We rasterised four constructions on 2^{22} pixels of $[0, 1]$. Each construction replaces every interval by its two end pieces:

1. self-similar, $r = 1/4$ (tier 1);
2. $r_k = 2^{-1/d_k}$ with $d_k = \frac{1}{4} + \frac{1}{2}\Omega_k$, using the BLC lower bounds as Ω_k (tier 1, by Proposition 3.1);
3. one random bit per stage selecting $r_k \in \{0.25, 0.10\}$, a stand-in for the bits of Ω (tier 2);
4. one random bit per interval (tier 3).

For each stage we record the LZMA-compressed size [16] minus that of an empty bitmap (1664 bits), which is an upper bound on K up to an additive constant, together with the box-counting dimension.

The results are shown in Fig. 1. The self-similar set saturates at about 10^3 bits for $L \gtrsim 8$. The tier-3 set grows fastest, reaching about 3×10^3 bits at stage 7 with 128 intervals, and the tier-2 set lies in between. LZMA can only approximate the true complexities from above: it cannot discover the $O(\log L)$ programs of tier-1 sets. The box-counting dimension of construction 2 approaches the Cesàro value of the stage ratios (0.29 at stage 6, against 0.31 from box counting). The limiting dimension $\frac{1}{4} + \frac{1}{2}\Omega_{\text{BLC}}$ is known only to lie in $[0.310, 0.75]$.

6 Discussion and physical scope

1. Theorem 1 uses a universal prefix-free machine only through the left-c.e. real Ω . Any left-c.e. real in $[0, 1]$ can be substituted, so the construction realises every left-c.e. dimension in $[\frac{1}{4}, \frac{3}{4})$. Extending it to other classes of reals, or to dimensions of sets defined by natural dynamics, is left open.
2. Proposition 3.1 shows that uncomputability of a geometric invariant is compatible with minimal complexity at every finite resolution. The two notions should not be conflated.
3. The constant $c_F \approx 2.4 \times 10^5$ bits is illustrative. Smaller ZFC-independent machines have since been reported, which only strengthens the argument of Section 4.
4. The numerical part is illustrative: LZMA gives upper bounds only, and the BLC computation gives a lower bound only, for a halting probability that we do not claim to be random.




Figure 1: Left: lower bound on the toy halting probability Ω_{BLC} against the maximal program length; the upper bound is trivial (see text). Middle: LZMA-compressed size above the empty-bitmap baseline (solid) and the complexity predicted by Propositions 3.1–3.3 up to $O(1)$ (dotted) for the four constructions; the horizontal line marks the illustrative Chaitin scale c_F . Right: running dimension per stage (lines) and box-counting estimates (crosses); the band is the interval for $\frac{1}{4} + \frac{1}{2}\Omega_{\text{BLC}}$ implied by the bounds on Ω_{BLC} .

Data and code availability

All numbers and Fig. 1 are produced by the Python script `chaitin_fractal.py`, available at <https://github.com/Ruqing1963/chaitin-fractal-complexity> and archived on Zenodo under DOI 10.5281/zenodo.23194419.

A Technical details

BLC enumeration and reduction. Terms of length exactly s with free indices below k are generated recursively. A variable with index $i < k$ has length $i + 2$, an abstraction has length $2 +$ that of its body (with $k \rightarrow k + 1$), and an application has length $2 +$ the sum of the lengths of its parts. Closed terms are those with $k = 0$. Reduction uses de Bruijn shifting and substitution with the leftmost-outermost (normal-order) strategy, which finds a normal form whenever one exists. A term is recorded as normalising if no redex remains within 400 steps; it is recorded as undecided if the step budget is exhausted or the term exceeds 3000 nodes.

Lower and upper bounds. The lower bound sums $2^{-|M|}$ over certified normalising terms. The upper bound adds the mass of undecided terms and the Kraft remainder $1 - \sum_{|M| \leq n} 2^{-|M|}$. Because every enumerated term is either certified or undecided, this upper bound equals 1 identically. A non-trivial upper bound requires the two ingredients described in Section 5.

Rasterisation and compression. Intervals are painted onto a uint8 array of 2^{22} cells using a difference array, then bit-packed and compressed with LZMA (preset 9, extreme). The empty-bitmap baseline of 1664 bits is subtracted. Construction stops once the smallest interval is shorter than one pixel.

Box counting. Occupied boxes are counted at dyadic sizes 2^{-j} , for $j = 3, \dots, 22$, restricted to box sizes not smaller than the current stage length. The dimension is the least-squares slope of $\ln N$ against $j \ln 2$.

References

- [1] M. Li, P. Vitányi, *An Introduction to Kolmogorov Complexity and Its Applications*, 4th ed., Springer (2019).
- [2] R. G. Downey, D. R. Hirschfeldt, *Algorithmic Randomness and Complexity*, Springer (2010).
- [3] G. J. Chaitin, “A theory of program size formally identical to information theory”, J. ACM 22 (1975) 329–340, doi:10.1145/321892.321894.
- [4] P. Martin-Löf, “The definition of random sequences”, Inform. Control 9 (1966) 602–619.
- [5] M. Braverman, M. Yampolsky, “Non-computable Julia sets”, J. Amer. Math. Soc. 19 (2006) 551–578, doi:10.1090/S0894-0347-05-00516-3.
- [6] T. S. Cubitt, D. Pérez-García, M. M. Wolf, “Undecidability of the spectral gap”, Nature 528 (2015) 207–211, doi:10.1038/nature16059, arXiv:1502.04135.
- [7] C. Moore, “Unpredictability and undecidability in dynamical systems”, Phys. Rev. Lett. 64 (1990) 2354.
- [8] G. J. Chaitin, “Information-theoretic limitations of formal systems”, J. ACM 21 (1974) 403–424.
- [9] A. Yedidia, S. Aaronson, “A relatively small Turing machine whose behavior is independent of set theory”, Complex Systems 25 (2016) 297, doi:10.25088/ComplexSystems.25.4.297, arXiv:1605.04343.
- [10] J. Tromp, “Binary lambda calculus and combinatory logic”, in *Randomness and Complexity: From Leibniz to Chaitin*, World Scientific (2007) 237–260, doi:10.1142/9789812770837_0014.
- [11] M. Stay, “Very simple Chaitin machines for concrete AIT”, Fundam. Inform. 68 (2005) 231–247, arXiv:cs/0508056.
- [12] K. Weihrauch, *Computable Analysis: An Introduction*, Springer (2000).
- [13] K. Falconer, *Fractal Geometry: Mathematical Foundations and Applications*, 2nd ed., Wiley, Chichester (2003), doi:10.1002/0470013850.
- [14] J. E. Hutchinson, “Fractals and self similarity”, Indiana Univ. Math. J. 30 (1981) 713, doi:10.1512/iumj.1981.30.30055.
- [15] OEIS Foundation, The On-Line Encyclopedia of Integer Sequences, sequence A114852, “The number of closed lambda calculus terms of size n ”, <https://oeis.org/A114852>.
- [16] J. Ziv, A. Lempel, “A universal algorithm for sequential data compression”, IEEE Trans. Inf. Theory 23 (1977) 337–343.